



CSC 355 Database Systems

Lecture 6

Eric J. Schwabe

School of Computing, DePaul University

Spring 2020



Topics:

- ◆ SQL queries
 - Review GROUP BY and HAVING
 - Query problems
 - Introduction to joins

Aggregate Functions

- ◆ Given an attribute, these functions take the values of that attribute in the set of returned rows and compute a single value from them
 - COUNT(...): Number of non-NULL values
 - SUM(...): Sum of the values
 - AVG(...): Average of the values
 - MIN(...): Smallest of the values
 - MAX(...): Largest of the values

GROUP BY

... GROUP BY *grouping attributes* ...

- ◆ Combines the rows into sets based on the value(s) of some attribute(s)
 - Can only display the value(s) of this attribute(s) and/or aggregate information for each group
 - If we group rows into sets, we cannot look at the values in the individual rows anymore...

HAVING

... HAVING *condition on groups* ...

- ◆ Includes only those groups that satisfy the condition
 - the condition may only involve the grouping attribute(s) and/or aggregate functions
 - can use all the same comparisons and logical operators as WHERE

Query Structure (again)

- ◆ General form of a query:

SELECT *list of expressions*

FROM *set of rows*

[WHERE *condition on rows*]

[GROUP BY *grouping attributes*]

[HAVING *condition on groups*]

[ORDER BY *ordering attributes*] ;

- ◆ Grouping goes after WHERE, before ORDER BY

Writing a Query

1. FROM: What table should I use?
2. WHERE: How do I indicate which rows to include?
3. GROUP BY: What attribute's values will define the sets? (May have to change SELECT * here...)
4. HAVING: How do I indicate which sets to include?
5. ORDER BY: How should I sort the rows/sets?
6. SELECT: What values do I have to compute and display?

Query Problems

- ◆ Find the number of workers in each department
 - (...whose salary is more than \$40,000)
- ◆ Find the average salary over the entire company
- ◆ For each department, find the salary of the highest-paid employee
- ◆ List the department names and their total budgets, ordered from the largest total budget to the smallest
- ◆ For each student, find the total number of classes they have enrolled in and the most recent year that the student enrolled in a class

Joins

- ◆ Data that is distributed among multiple tables can be combined into a single set of rows for use in a query using different types of *joins*:
 - Inner joins (equi-join, natural join)
 - Outer joins (left, right, full)

Cartesian Product

- ◆ What if we list two tables in the FROM?
- ◆ The rows in the result come from combining all pairs of rows from the two tables – the *Cartesian Product* of the tables
 - (This is sometimes called the “cross join”...)
- ◆ This is almost certainly more rows than we want – most combinations are meaningless!

Equi-Join

- ◆ An equi-join keeps only those rows where the two combined rows agree on the shared attribute(s):

```
...FROM TABLE1, TABLE2  
WHERE  
TABLE1.Attribute = TABLE2.Attribute;
```

Natural Join

- ◆ Like an equi-join, but one of the duplicated columns is removed (the most common join):

SELECT *all but the duplicated attribute(s)*
FROM *TABLE1, TABLE2*
WHERE
TABLE1.Attribute = TABLE2.Attribute;

Inner Joins

- ◆ These are both examples of *inner joins*
- ◆ In an inner join, the Cartesian Product is restricted to only include the combined rows that satisfy some condition
 - condition is usually equality in some shared key
 - e.g., equi-joins, natural joins

Inner Joins

- ◆ Rather than list of tables in the FROM and a WHERE condition, can use:

FROM *TABLE1* INNER JOIN *TABLE2*
ON *condition*

Join Example

COURSES(CourseNumber, CourseName)

SECTIONS(SectionID, CourseNumber, SectionNumber)

ENROLLMENTS(StudentID, SectionID)

STUDENTS(StudentID, FirstName, LastName)



Next:

- ◆ More SQL Queries
 - Inner joins
 - Outer joins
 - Query examples
 - Set operations