

CSC 355 Database Systems

Lecture 12

Eric J. Schwabe

School of Computing, DePaul University

Spring 2020

Today:

- ◆ Relational Database Design
 - BCNF decomposition algorithm
 - Dependency preservation
 - Lossless join

Relational Database Design

- ◆ Starting with $R(A_1, A_2, \dots, A_n)$, the decomposition $D = \{R_1, R_2, \dots, R_m\}$ should satisfy the following conditions:
 1. The union of the R_i 's is R
 2. Redundancy has been removed from relations in D
 3. The functional dependencies in R are preserved
 4. The original relation R can be recovered from D
- ◆ We are formalizing conditions 2.-4. ...

Functional Dependencies

- ◆ A set of attributes $Y = \{Y_1, Y_2, \dots, Y_n\}$ is *functionally dependent* on a set of attributes $X = \{X_1, X_2, \dots, X_m\}$ if and only if every pair of tuples that have the same values for X must also have the same values for Y
 - Also “ X functionally determines Y ” or “ $X \rightarrow Y$ ”
 - X is called the *determinant*
- ◆ (Less formally, “the values of X uniquely determine the values of Y ”...)

Keys

- ◆ A set of attributes X is a *superkey* of R if X determines all attributes of R
- ◆ A set of attributes X is a *candidate key* of R if X is a superkey, but no proper subset Y of X is a superkey
- ◆ An attribute is *prime* if it is contained in some candidate key (and is *non-prime* otherwise)

BCNF

- ◆ A relation R is in Boyce-Codd Normal Form (BCNF) if for every non-trivial functional dependency $X \rightarrow Y$ in R , X is a superkey
 - “Every determinant must contain a candidate key”
 - A relation in BCNF will not have any redundancy, since every functional dependency in the relation will have a superkey as its determinant

BCNF Decomposition Algorithm

Set $D = \{R\}$

While there is some Q in D that is not in BCNF:

 Choose a Q that is not in BCNF

 Find an $X \rightarrow Y$ in Q that violates BCNF

 Replace Q with two relations:

$Q - Y$ and $(X \text{ union } Y)$

(When algorithm is done, all relations will be in BCNF)

Projections

- ◆ Suppose we have a set of functional dependencies F in the relation R
- ◆ For a relation R_i , consider all $X \rightarrow Y$ that can be derived from F where both X and Y are subsets of R_i
- ◆ This set is called the *projection of F on R_i*
 - It represents the set of all constraints that F puts on the attributes of R_i

Dependency Preservation Property

3. The union over all i in $\{1, \dots, m\}$ of the projections of F on R_i is equivalent to F
 - ◆ We want the set of all projections to be equivalent to F – that is, the decomposition neither destroys any functional dependencies in F nor introduces any new ones...
 - Not all decompositions have this property!

Restrictions

- ◆ The *restriction* of a relation state r to a set S is the set of distinct tuples obtained from r by including only the values of the attributes in S from each tuple
- ◆ Restrictions that overlap can be combined using a natural join on the attributes they share
 - Ideally, this will yield a result that is still a restriction of the original relation state r ...

Lossless Join Property

4. For every relation state r of R , the natural join of the restrictions of r to the relations R_1, R_2, \dots, R_m in the decomposition the same as r
- That is, if we take the restrictions of any relation state and join them back together, we will get the original relation state – no *spurious tuples* are added
 - Not all decompositions have this property!

Relational Database Design

- ◆ Starting with $R(A_1, A_2, \dots, A_n)$, the decomposition $D = \{R_1, R_2, \dots, R_m\}$ should satisfy the following conditions:
 1. The union of the R_i 's is R
 2. Each of the R_i 's is in BCNF
 3. D has the dependency preservation property
 4. D has the lossless join property

Binary Lossless Join Test

- ◆ $D = \{R_1, R_2\}$ has the lossless join property if and only if one or both of the following hold:
 1. $(R_1 \cap R_2) \rightarrow (R_1 - R_2)$ can be derived from F
 2. $(R_1 \cap R_2) \rightarrow (R_2 - R_1)$ can be derived from F
- That is, if and only if the intersection between the two sets of attributes is a superkey in one of the relations...

General Test for Lossless Join

1. Create a matrix S with a row i for each R_i and a column for j for each A_j
2. Set each $S(i,j) = "b_{ij}"$
3. For each entry (i,j)
If relation R_i includes A_j , then set $S(i,j) = "a_j"$
4. Repeat the following loop until there are no changes to S :
For each $X \rightarrow Y$ in F
For all rows in S that have the same symbols in all columns in X ,
set all of the columns in Y in those rows to agree as follows:
If any row has a_j for the columns, set all rows to that same a_j
If not, choose one of the b_{ij} s and set all rows to that same b_{ij}
5. If any row has all a_j 's, return true (D has the lossless join property);
if not, return false (D does not have lossless join property).

Next:

- ◆ Finish lossless join
- ◆ Third Normal Form (3NF)
- ◆ Minimal basis
- ◆ Algorithm for 3NF decomposition