



CSC 355 Database Systems

Lecture 10

Eric J. Schwabe
School of Computing, DePaul University
Spring 2020



Today:

- ◆ Finish Transactions
- ◆ Midterm Exam Information
- ◆ Introduction to Relational Database Design

Transactions

- ◆ A *transaction* is a collection of SQL statements that must be executed as a unit
- ◆ Transactions have “*ACID*” properties:
 - Atomicity: Execute completely or not at all
 - Consistency: Satisfy all database constraints
 - Isolation: Execute “separately” from others
 - Durability: Once completed, results are permanent

Serializable Isolation

- ◆ Transactions must behave as though they were run serially (first one, then the other)
- ◆ Usually implemented by “locking” the tables (or parts of tables) used by a transaction
 - Other transactions using the same tables will have to wait until they are released
 - Other transactions using other tables could run at the same time

Read Committed Isolation

- ◆ Transaction operations can be interleaved
- ◆ If one transaction tries to read data that were written by another, it can only see the changes that have been committed
 - Can't be rolled back, but could be changed later
 - Multiple queries of the same table might not yield the same results... “non-repeatable reads”, including “phantoms”...

Read Uncommitted Isolation

- ◆ Transaction operations can be interleaved
- ◆ If one transaction tries to read data that were written by another, it can see all changes, even if they have not been committed
 - Could be rolled back by the other transaction!
 - Transaction might make decisions based on values that are later rolled back and so were never really there ... “dirty reads”...

Isolation Levels

- ◆ **SERIALIZABLE:** Transactions must appear to run serially – cannot read any changes from others
- ◆ **[REPEATABLE READ:** Can read committed changes that only add data (allows only phantoms)]
- ◆ **READ COMMITTED:** Can read all committed changes (allows all non-repeatable reads)
- ◆ **READ UNCOMMITTED:** Can read all changes (allows all non-repeatable reads and dirty reads)

Transactions in Oracle

- ◆ SET TRANSACTION to start a transaction
 - Specify transaction NAME
 - Can also specify ISOLATION LEVEL
 - READ COMMITTED (default) or SERIALIZABLE
 - COMMIT or ROLLBACK ends transaction
 - SQL statements that modify data start a transaction implicitly... COMMIT to end it

Midterm Exam Information

- ◆ Exam Monday 5/4, at regular class time (90 minutes)
- ◆ Exam will be given as a quiz in d2l
- ◆ Lecture slides will be available, but no other sources of information (electronic, printed, or human) may be consulted
 - I will have you sign a statement agreeing to this as part of exam
- ◆ Sections 1.1-1.3, 2.1-2.3, 6.1-6.6, 7.1-7.3 covered
- ◆ Multiple choice, short answers, writing and/or evaluating SQL queries and transactions
- ◆ Review outline has been posted
- ◆ Optional Q&A session Friday 5/1, via Zoom

Relational Database Design

- ◆ Start with a set of attributes

$$R = \{A_1, A_2, \dots, A_n\}$$

- (Can also be written as a *universal relation*
 $R(A_1, A_2, \dots, A_n) \dots$)

- ◆ Construct a *decomposition* of R into relations

$$D = \{R_1, R_2, \dots, R_m\}$$

- Each R_i is a subset of R

Relational Database Design

- ◆ The decomposition $D = \{R_1, R_2, \dots, R_m\}$ should satisfy the following conditions:
 1. The union of the R_i 's is R
 2. Redundancy has been removed from the R_i 's
 3. Dependencies among attributes in R are preserved
 4. The original relation R can be recovered from D
- ◆ Conditions 2.-4. have to be formalized...

Redundancy

- ◆ *Redundancy* occurs when more than one record in a table stores the same information
 - Wastes space
 - Allows *update* and *deletion* anomalies
- ◆ Remove redundancy by identifying (and removing) *functional dependencies* in R

Functional Dependencies

- ◆ A set of attributes $Y = \{Y_1, Y_2, \dots, Y_n\}$ is *functionally dependent* on a set of attributes $X = \{X_1, X_2, \dots, X_m\}$ if and only if every pair of tuples that have the same values for X must also have the same values for Y
 - Also “ X functionally determines Y ” or “ $X \rightarrow Y$ ”
 - X is called the *determinant*
- ◆ (Less formally, “the values of X uniquely determine the values of Y ”...)

Functional Dependencies

- ◆ “Every pair of tuples that have the same values on X also have the same values on Y”
 - For X to functionally determine Y, this condition must be *satisfied by every possible relation state*
 - If *some relation state does not satisfy* the condition because two tuples have the same values on X but different values on Y, then X does not functionally determine Y

Finding Functional Dependencies

- ◆ DVD (DVDID , MovieID , Title , Genre , Length , Rating)
- ◆ GRADING (CNumber , CTitle , SID , SName , Grade)
- ◆ (We do not typically include *trivial* functional dependencies, where Y is a subset of X...)

Okay ... but why?

- ◆ Redundancy comes from functional dependencies whose determinants do not include a complete candidate key of R
- ◆ We use the functional dependencies to construct decompositions of R
- ◆ How do we measure the quality of the resulting decompositions?



Next:

- ◆ Midterm exam Monday 5/4
 - Q&A session Friday 5/1
- ◆ Next lecture will be posted Wednesday 5/6